

Return-Weighted Risk for Navigating an Evolving AI Landscape

Reconnecting Risk with Value in AI Risk Management

Honggao Cao

August 2026

Abstract

Organizations undertake AI activities to create value, yet often have to decide whether to continue or expand them while still learning what value they deliver and what risks they introduce. Changes in capabilities, configurations, users, costs, and alternatives can weaken the evidence supporting an earlier approval. We propose return-weighted risk (RWR) as an activity-level governance approach that connects an evolving value case with a risk case subject to appropriately independent assessment and challenge at each authorization decision. Three sequential tests organize that decision: whether the action satisfies applicable constraints, whether residual risk is permissible, and whether the current value case justifies taking that permissible risk. Intended value directs additional risk inquiry while a baseline assessment covers risks beyond the intended benefit. Feasible alternatives remain part of the comparison, and material change in purpose, evidence, or operating conditions can reopen authorization. A hypothetical small-business-lending example shows why a successful pilot may support further controlled learning without yet supporting expansion. The approach makes the basis for continuing action explicit; its effects on governance decisions and outcomes require empirical evaluation.

Keywords: AI governance; AI risk management; return-weighted risk; activity-level governance; concurrent learning; continuing authorization; material change; value realization

1. The Governance Problem in an Evolving AI Environment

Organizations adopt artificial intelligence because they expect beneficial change—financial, operational, customer-related, workforce-related, strategic, scientific, public, or informational. The same activities can also create or amplify risks involving reliability, privacy, security, fairness, compliance, model behavior, third-party dependence, operational resilience, and human reliance. What makes the present environment especially difficult is that organizations must often learn about value and risk while the activity is already being developed, piloted, deployed, and changed.

This creates a moving governance target. AI capabilities, deployment configurations, user behavior, available alternatives, and organizational objectives can shift before the next decision point. A sound assessment today may therefore lose decision relevance sooner than governance processes built around more stable assumptions would expect.

When both value and risk are unsettled, an organization has to decide without treating either assessment as final. It may still be learning how a task improvement becomes organizational value, how much that improvement costs, and what exposure develops in use. Waiting for complete knowledge can prevent

useful learning; treating the initial business and risk cases as settled can allow authorization to persist after its justification has changed. The practical difficulty becomes clearer in an expansion decision.

1.1 A motivating decision: Aster Vale

Consider Aster Vale Regional Bank, a hypothetical regional bank that completed a six-month pilot of Navigator, a vendor-developed generative-AI system used to organize documents and draft portions of credit memoranda for small-business loans. Navigator does not approve or decline applications, assign credit ratings, determine prices, or communicate directly with applicants. Underwriters retain formal decision authority and certify every completed memorandum.

The pilot appears successful from one perspective. Navigator reduces underwriter preparation time by 28 percent, experienced underwriters report spending less time sorting documents, and the system remains within its restricted scope. The bank identifies no autonomous credit decision, serious privacy incident, or explicit breach of approved pilot risk limits.

The broader evidence is mixed. Application-to-decision time declines only 7 percent because document preparation is not the principal source of delay. Repeated requests for applicant information fall only slightly, abandonment scarcely changes, and the anticipated improvement in credit access is not clearly established. A non-AI workflow redesign produces a similar improvement in end-to-end processing time.

Navigator performs best with standardized digital information; scanned, handwritten, multilingual, or unconventional records require more manual work. Some newer underwriters increasingly begin with Navigator summaries, and source accessibility deteriorates. The bank gains useful knowledge about generative AI while becoming more dependent on a vendor whose system it understands only partially.

Management nevertheless proposes expansion to all small-business lending teams and larger loans. Its revised rationale emphasizes processing capacity, avoided hiring, and strategic AI capability. The pilot, however, had originally been justified mainly by faster decisions, lower applicant burden, and improved access. The revised rationale may be legitimate, but it is materially different.

The bank therefore faces a specific question: given what it has learned about Navigator's value and risk, does the evidence support broader deployment? The question is clear even though the answer requires judgment. Success within the pilot's restricted scope establishes a starting point for that judgment, rather than settling it.

1.2 Return-weighted risk as a governance approach

The Aster Vale example raises the question at the center of this paper: How should organizations govern AI activities when they must learn about value realization and an evolving risk profile at the same time as they decide whether to continue, adapt, expand, restrict, suspend, or terminate those activities?

We propose return-weighted risk (RWR) to organize this continuing decision. A value case states why an activity is worth undertaking; a risk case assesses its possible adverse effects and the remaining exposure after controls. We keep those cases distinct and bring them together for a specified action. Intended value helps direct additional risk inquiry, acceptable risk makes the action eligible for value judgment, and

feasible alternatives help establish whether the proposed exposure is justified. Changes in either case can require a new authorization decision.

We first examine why evidence about an evolving AI activity may lose relevance to the next decision (Section 2), then draw on established approaches to risk, value realization, and learning (Section 3). Those foundations lead to the two cases and three authorization tests in Sections 4 and 5. Section 6 explains how an authorized learning stage can update both cases. We then apply the sequence to Aster Vale and consider its practical use and empirical evaluation before concluding.

2. Why AI Makes the Problem More Acute

Aster Vale’s expansion request changes both what the bank hopes to achieve and the conditions under which Navigator would operate. To understand why the pilot cannot settle that request, we need to examine what can change between an assessment and the next decision. Time, deployment conditions, and the link from task outcomes to organizational value each affect the evidence the bank can use.

2.1 Rapid change makes point-in-time assessment transient

AI changes not only what organizations know, but how long that knowledge remains decision-relevant. A point-in-time assessment can be carefully performed and still become stale as capabilities expand, vendors update systems, deployment patterns shift, costs change, or feasible alternatives improve. In some cases, the activity being assessed is no longer quite the same activity by the time the next decision is made.

The problem affects both sides of the decision. Expected value can change as capabilities and alternatives shift, while new evidence can change what the organization believes has actually been realized. Risk can also change as configurations, users, and operating conditions create different exposures. Continuing governance therefore cannot assume that a sound initial assessment remains stable simply because it was sound when made.

In April 2026, U.S. banking agencies explicitly excluded generative and agentic AI from revised model-risk-management guidance because those technologies were considered “novel and rapidly evolving” ([1]). The agencies also announced plans for further information gathering focused on banks’ use of AI, including generative and agentic systems.

The exclusion leaves banks responsible for determining appropriate governance and controls through their risk-management practices ([1]). We use the episode as an illustration of the changing governance setting. The continuing-authorization problem also arises within established frameworks: the organization must decide which parts of an earlier assessment remain applicable to the activity it now proposes.

2.2 Purpose and the deployed activity can change

A formal change in scope is usually easy to see: an AI system may move from document preparation to recommendation, from a restricted user group to an enterprise workflow, or from an internal aid to customer-facing interaction. More difficult is the gradual change that occurs when users begin relying on the system differently, prompts or retrieval change, the vendor updates the product, or an internal aid

begins shaping decisions more heavily than originally intended. Governance labels can remain stable while the effective use changes underneath them.

These changes are consequential because the deployed activity includes the arrangements through which a model is used. Data, retrieval, prompts, tools, human review, workflow, vendor changes, controls, and operating context all influence what happens in use. Aster Vale, for example, is not governing Navigator in isolation from its underwriters, source access, document types, verification practices, and workflow. Evidence from a pilot therefore describes a configured activity under particular conditions; it does not automatically establish how the same underlying model will perform when those conditions change.

2.3 From outcome to value

Changes in those arrangements also affect whether a favorable task result advances the intended purpose. Navigator can reduce document-preparation time without producing a comparable improvement in end-to-end lending performance. We distinguish the observed outcome—less preparation time—from the value that outcome creates through faster decisions, reduced applicant burden, or another relevant objective. The link between the two has to be supported rather than assumed.

Even when the downstream outcome improves, a further question remains: which changes produced the improvement? AI use, training, workflow redesign, user selection, policy changes, and their interactions may all contribute. OECD evidence that many firms struggle to estimate AI return on investment is consistent with this measurement difficulty ([2]). For governance, evidence about the configured activity can support a decision even when AI's separate contribution remains uncertain. A claim that AI itself created the improvement requires a stronger basis for attribution, especially when comparing the activity with a non-AI alternative.

2.4 Risk can shift across users and conditions

The value pathway can also work differently across users and cases, creating different risk exposures. Experienced employees may save time while newer employees become more reliant on AI summaries; standardized documents may work well while nonstandard records require more verification or produce poorer results. An average gain can conceal these differences. Both cases therefore need to identify who benefits, who bears exposure, and under which inputs and controls those effects arise.

2.5 Evidence may not travel

These differences become especially important when the organization proposes expansion. A pilot supplies evidence about its users, document types, controls, and scale. Expansion may change each of those conditions, along with incentives and consequences. Evidence from experienced underwriters working with standardized documents under close oversight, for example, may leave important questions unanswered about less-experienced users handling more varied documents with reduced quality assurance. The organization should identify which findings can reasonably carry forward and which proposed conditions require additional evidence.

2.6 Governing while learning

The preceding questions leave the organization with incomplete knowledge on both sides of the decision. Some uncertainty concerns whether the intended value will emerge; some concerns exposure under changed operating conditions. A smaller action may help answer those questions, provided that the action used to learn is itself justified.

The situation resembles crossing a river by testing the footing ahead. The destination may be understood even when the complete path is not, and each controlled step supplies information for the next. In governance terms, each step must remain within acceptable exposure, and the information expected from it must bear on a subsequent decision. This is the learning logic we develop for RWR in Section 6.

The need is therefore to connect current evidence to a specified action, learn within authorized limits, and reconsider the action when its justification changes. Several established approaches already supply important parts of that logic.

3. Foundations for the RWR Approach

We turn to those approaches because each answers a question raised by Aster Vale. Risk governance connects exposure to objectives and limits. Benefits realization asks whether the expected value actually emerges. Adaptive approaches explain how action can produce evidence for later decisions. Benefit-risk traditions clarify the constraints that expected value cannot override. Together, they provide the foundations for RWR’s continuing-authorization decision.

3.1 Risk, value, and organizational purpose

Risk management has long been connected to organizational objectives rather than treated as an isolated control function. The Committee of Sponsoring Organizations of the Treadway Commission (COSO) enterprise risk management framework integrates risk with strategy and performance ([3]). In banking, Office of the Comptroller of the Currency (OCC) new-activity guidance explicitly connects strategy, risks and returns, alternatives, performance criteria, exit strategies, and ongoing review ([4]). Revised model-risk and third-party guidance adds expectations for monitoring, validation, accountability, and continuing fitness for purpose ([1], [5]).

AI governance follows much the same direction. The National Institute of Standards and Technology (NIST) Artificial Intelligence Risk Management Framework (AI RMF) asks organizations to consider intended purpose, beneficial uses, assumptions, limitations, and positive and negative impacts while managing risk across the lifecycle ([6]). Organizational AI-governance research similarly connects AI use with organizational strategy, objectives, values, and obligations ([7]).

The field is still developing. Reviews identify fragmentation and implementation gaps in organizational AI governance, while design-oriented work has proposed more practice-oriented requirements ([8], [9]). Financial-sector surveys report substantial expectations for AI alongside concerns about data, privacy, security, bias, model complexity, third-party dependence, and skills ([10], [11]). OECD evidence adds a practical difficulty: firms often struggle to estimate AI return on investment, and experimentation does

not guarantee success ([2]). Taken together, this work shows that risk can be considered alongside the organizational purpose for undertaking an activity, even when the expected value remains difficult to establish.

3.2 Value realization

The connection between risk and purpose still leaves a question about delivery: did the approved activity create the value used to justify it? Benefits realization management addresses that question by specifying expected benefits, assigning owners, measuring against baselines, following benefits beyond project delivery, and reconciling them with the business case ([12], [13]). Responsibilities and resources can break down in practice ([14]). Aster Vale illustrates why follow-through matters: Navigator produced a clear task-level improvement while the broader customer and access value remained uncertain.

3.3 Learning under uncertainty

Benefits realization does not eliminate the problem that some uncertainty can only be reduced through experience. Adaptive management addresses that problem more directly: stakeholders, objectives, alternatives, models, monitoring, and observed responses are used to update decisions as evidence emerges ([15]). Later work recognizes that learning can change not only technical understanding but also objectives or institutional decision assumptions ([16], [17]). Dynamic Adaptive Policy Pathways and real-options reasoning add related ideas of staged commitment, adaptation triggers, and options to expand or abandon as information improves ([18], [19]).

Incomplete knowledge, then, does not always require either full commitment or complete delay. Limited action can sometimes be used to learn, provided that the exposure required to learn remains acceptable.

3.4 Constraints on value-risk trade-offs

Expected value cannot justify every level or type of risk. Existing benefit-risk approaches already recognize this limit. Medical-device regulation can consider uncertainty while maintaining safety and effectiveness requirements ([20]). The Health and Safety Executive (HSE) uses a tolerability-of-risk framework alongside a continuing duty to reduce risk so far as is reasonably practicable ([21]). Multicriteria benefit-risk research warns that simple aggregation can produce counterintuitive recommendations ([22]). Responsible-innovation scholarship adds anticipation, reflexivity, inclusion, and responsiveness, extending attention beyond organizational return alone ([23]). Together, these traditions support considering constraints and risk limits before asking whether the remaining exposure is justified by value.

3.5 The focus of RWR

These traditions establish the relationships we carry into RWR: purpose informs the activity, benefits require evidence after approval, learning can support staged action, and constraints limit what value may justify. The remaining task is to bring those relationships to bear on a particular continuation or expansion decision as both cases change.

These foundations explain the components RWR uses. The closest remaining question is how RWR differs from AI-governance approaches that already connect value creation and risk mitigation. Recent work frames that relationship as a persistent organizational tension and develops the Paradox-based Responsible AI Governance (PRAIG) framework to guide governance strategies across different contexts and AI portfolios ([24]). We focus here on the authorization of a specified activity: whether its current value and risk cases support the proposed action.

Consider Aster Vale again. The pilot produced favorable outcomes, but the original value rationale is only partly supported. Management now proposes a different rationale and broader exposure, while a non-AI alternative remains feasible. The question for RWR is whether, taken together, the current value case and independently assessed risk case still support the same action.

That continuing authorization decision is the focus of RWR. Table 1 summarizes how this focus relates to several adjacent governance perspectives.

Table 1. Governance perspectives on the Aster Vale expansion decision

Perspective	Primary question in the Aster Vale decision
Risk management	Is residual risk acceptable for the proposed scope and controls?
Benefits realization	Are intended benefits defined, owned, measured, and being realized?
Adaptive management	What should be learned next, and how should action adapt as evidence changes?
PRAIG	Which governance strategies address the continuing tension between AI value creation and risk mitigation?
RWR	Do the current value and risk cases, relative to feasible alternatives, support the proposed action?

The perspectives in Table 1 are complementary; established guidance such as the OCC’s new-activity principles can support a similar disposition ([4]). RWR brings their evidence to a common decision point. To make that connection usable, we next specify what belongs in each case before examining how authorization joins them.

4. Return-Weighted Risk: The Governance Approach

4.1 RWR in plain language

An organization approves an AI activity for a reason and under particular conditions. As the activity operates, evidence may change both the reason for undertaking it and the exposure it creates. We use RWR to keep the value case connected to an independently assessed risk case and ask whether both support the proposed action under current conditions. The unit of analysis is the deployed activity—its AI system, users, workflow, data, controls, and scope—rather than the model in isolation.

4.2 Why “Return-Weighted Risk”?

The name expresses the direction of inquiry. Risk-adjusted return reasoning asks how risk affects the attractiveness of an expected return. We also ask the reverse question: given the value an AI activity is intended to create, which dependencies and possible adverse effects deserve additional investigation? For Aster Vale, an intended improvement in access directs attention to applicants with nonstandard documents; an intended reduction in staffing needs directs attention to verification capacity and workload. Those questions arise from the value case before they can be answered by the risk assessment.

“Weighted” therefore describes an emphasis in inquiry and judgment. RWR weights the direction of inquiry, not the risk conclusion. It retains independent risk conclusions and uses three authorization tests rather than a mathematical weighting formula or common value-risk score. We use return in the name to preserve its connection with investment reasoning. The operational object is a value case, because relevant value may include financial results, service, workforce effects, access, strategic capability, or learning that cannot be reduced to one monetary measure.

4.3 The value case

The value case makes the reason for undertaking the activity explicit. Work may become faster, service may improve, costs may fall, or a useful capability may be developed. The case explains how those outcomes are expected to advance an organizational or stakeholder objective. In Aster Vale’s original case, less preparation time was an intermediate outcome; faster decisions and reduced applicant burden were the intended downstream benefits.

The value case therefore keeps enough information to understand what the organization is trying to achieve, for whom, relative to what alternative, at what cost, and with what evidence. Relevant details will vary with the decision, but usually include the important outcomes, stakeholders, resources and costs, alternatives, and supporting evidence. What the organization intends to create is not necessarily what it can reasonably expect, and neither is the same as what it later observes or can support as realized value. As evidence changes, expectations should change with it. Attribution requires similar care: improvement under an AI-enabled activity does not by itself establish how much of that improvement was caused by AI rather than workflow, training, user selection, or other complementary changes.

For continuation, the comparison is forward-looking. Realized benefits inform what may happen next, but prior spending and past success do not by themselves justify further exposure. Remaining costs, usable capabilities already created, switching consequences, and feasible alternatives affect the value of the next action.

4.4 The risk case

The risk case examines the activity whose value has just been described: what can go wrong, who may be affected, how serious the possible effects are, and whether controls make the remaining exposure acceptable. We use risk here to mean possible adverse effects under uncertainty, and residual risk to mean the exposure remaining after the proposed controls. The assessment applies to the specified users, scope, and conditions; a permissible pilot risk case may leave the expansion risk case unresolved.

Expected spending on technology, people, verification, controls, and other resources belongs in the value case. Uncertainty about spending or resource availability can also create risk, but counting an expected cost as risk would blur the two judgments. Controls connect the cases: additional verification may reduce exposure while consuming some of the anticipated time saving. The risk case assesses the remaining exposure, and the value case includes the resources needed to achieve that level of control.

A baseline risk assessment supplies coverage beyond the intended benefit, including reliability, privacy, security, fairness, third-party dependence, and other relevant adverse effects. Purpose-directed inquiry adds questions to that coverage. In Aster Vale, the access objective makes performance on nonstandard documents especially relevant, but it gives the bank no reason to reduce privacy or security inquiry. The emphasis changes what the organization investigates, while the evidence and applicable limits determine the risk conclusion.

4.5 Coupled governance

The two cases meet at one authorization decision because the organization undertakes the same activity to obtain value and accepts exposure in doing so. Business sponsors support the value case; risk assessment and challenge establish the risk case with independence appropriate to the organization. The authorization body decides whether the proposed action is justified using both. Evidence may move between the cases, but a sponsor's expected benefit cannot determine whether a risk exists, how severe it is, or whether it falls within authorized limits.

5. Continuing Authorization Under RWR

Having separate cases makes the decision clearer, but it does not yet tell us how to proceed when one case is favorable and the other is weak. We organize continuing authorization through three tests. Their order distinguishes an action that may proceed from one that is worth undertaking under the proposed conditions.

5.1 Three tests

We apply the tests to a specified action, including its scope, controls, and conditions:

Admissibility asks whether the action satisfies applicable legal, ethical, and policy constraints, including constraints the organization treats as non-compensable. Non-compensable means that an expected benefit cannot offset failure to satisfy the constraint.

Risk permissibility asks whether independently assessed residual risk remains within authorized limits for the proposed conditions. Value justification then asks whether the current value case justifies undertaking that permissible exposure.

The order matters. Some actions should not proceed regardless of expected value; they fail the admissibility test. Other actions may be admissible but still carry more residual risk than the organization is prepared to accept. Only after an action passes those two tests does RWR ask whether the value case is strong enough to justify taking the permissible risk. Passing the risk test therefore does not mean approve; it means the activity is eligible for the value-justification decision.

The tests are sequential in judgment, although evidence gathering may proceed concurrently. Table 2 summarizes the order. A test that remains unresolved has not been passed for the proposed action; the organization must obtain adequate evidence, revise the action, or withhold authorization. This distinguishes lack of a basis to approve from an affirmative finding that the action is prohibited.

Table 2. RWR three-test authorization sequence

Test	Decision question	If not satisfied
Admissibility	May this action proceed under the applicable constraints?	Withhold authorization.
Risk permissibility	Is residual risk within permitted limits for the proposed conditions?	Withhold authorization in the proposed form.
Value justification	Does the current value case justify this permissible exposure?	Withhold authorization; adapt, narrow, or reconsider.

5.2 Value justification as governed judgment

Once the first two tests are satisfied, the decision turns to value justification. Some effects can be expressed in money; others concern customer experience, workforce effects, access, strategic capability, or learning. We therefore treat value justification as a governed judgment: the appropriate authority evaluates the evidence, alternatives, costs, and uncertainty on a record open to challenge and accountability. The record should explain which benefits carry the decision and how competing stakeholder interests were considered, even when those interests cannot be placed on a common scale.

Evidence adequacy follows the proposed action. A restricted pilot may be justified by a credible learning objective and limited exposure, while enterprise expansion requires evidence applicable to the broader use. A favorable finding for one action cannot simply be carried over to the other. This makes the strength of the value case and the permitted scope of authorization mutually relevant without changing the independent risk conclusion.

Passing the tests authorizes the specified action subject to its recorded conditions; it does not establish that the action is the best feasible choice. The next question is comparative: could the justified purpose be pursued through a different activity with a stronger overall case?

5.3 Avoiding unnecessary exposure

Passing all three tests can leave several feasible options. The organization then has to explain why it would choose the proposed activity over those alternatives. We ask whether a feasible option can pursue the justified purpose with less residual risk, taking account of differences in value, costs, controls, and implementation. For Aster Vale, workflow redesign's similar end-to-end gain makes broader Navigator exposure a question requiring justification. Less AI-related exposure alone does not establish that redesign has lower overall risk or greater value; its own effects and resource requirements belong in the comparison.

5.4 Alternatives remain live

That comparison remains relevant after approval because alternatives can improve or become newly feasible. Costs may shift, another tool may become available, or experience may reveal a different bottleneck. If an alternative can deliver comparable value with lower cost or exposure, the authorization body should reassess the reason for retaining the current activity. Transition costs, disruption, other stakeholder effects, and the alternative's own risk can still support continuation. The decision depends on the full comparison, rather than the existence of some value from the current use.

The comparison need not rebuild the full business and risk case for every alternative at every review. It should be refreshed when decision-relevant evidence changes materially. Navigator's preparation-time improvement, for example, looks different once Aster Vale learns that workflow redesign produces almost the same end-to-end improvement.

5.5 Material change and reauthorization

A change is material for RWR when it could plausibly change the outcome of one of the three tests. Revised purpose, new users, greater scale, weaker controls, different alternatives, or new evidence can each do this. Quantitative thresholds can identify some changes; others require judgment. Scheduled reconsideration can also be useful where uncertainty is high. The practical task is to specify which changes must reach the authorization body and how the activity will be governed while review is pending. A material change reopens the decision because the organization may be accepting exposure for a different reason or under different conditions.


For example, moving a document-preparation AI into credit recommendations changes the activity even if the vendor and model remain the same. The earlier authorization then supplies evidence for reconsideration, rather than authority for the new use. If evidence for expansion remains inadequate, the organization can ask whether a narrower learning stage would meet the tests.

6. Governing Concurrent Learning

6.1 The operating cycle

A narrower stage can answer a decision question that the proposed expansion cannot yet satisfy. Its authorization, however, has to stand on its own: the learning action must be admissible, its residual risk must be permissible, and the expected value of the information must justify its costs and exposure. Unresolved risk permissibility for the learning action itself is a reason to withhold that authorization. When all three tests are met, the organization can use a controlled stage to update the cases for its next decision. Figure 1 shows how the authorization tests enter that cycle.

Figure 1. RWR operating cycle



The organization specifies the purpose and the question to be learned, evaluates both cases for the proposed stage, and authorizes action only if the three tests are satisfied. Observations then update the cases and may support continuation, adaptation, expansion, restriction, suspension, or termination. Material change can reopen authorization before the planned stage ends. The cycle's return path represents reconsideration, not automatic permission to repeat the activity.

6.2 What needs to be learned

Learning has value in RWR only when it can inform a decision. Before beginning a learning stage, the organization should be clear about what it needs to know and what it would do differently depending on the answer. That usually means identifying the question being tested, what will be observed, what comparison matters, and what result would change the next authorization decision.

Aster Vale, for example, does not need another pilot simply to accumulate more experience with Navigator. It needs evidence that bears on specific decisions: whether source access can be restored, whether less-experienced users can use the system without excessive reliance, and whether Navigator produces enough incremental value over workflow redesign to justify broader exposure. More information is useful only when it can change what the organization does.

The learning plan should also identify who will review the findings, when the authorization will be reconsidered, and which observations require earlier restriction or suspension. These arrangements connect data collection to action. A pilot that produces information without a decision pathway can become a continuing activity whose learning rationale is never reassessed.

6.3 Shared evidence, distinct judgments

The observations collected during a learning stage can bear on both cases. Verification time, for example, affects the achievable time saving and may also indicate difficulty checking source material. The value assessment asks how that burden changes the benefit and cost of the activity; the risk assessment asks what it means for reliable review. Independence concerns the judgment and ability to challenge it, rather than separate datasets or isolation from business information.

The judgments can therefore inform one another without substituting for one another. A value claim may identify a dependency requiring further risk inquiry; a risk finding may require controls whose cost changes the value case. Expected benefit cannot turn an unsupported risk conclusion into a favorable one, and effective controls cannot establish that the activity is worth its cost. Authorization brings the two judgments together after their respective questions have been answered.

6.4 Decision authority and challenge

RWR can operate within an existing organizational structure. The value case needs an owner, the risk case needs credible independent assessment and challenge, and someone with appropriate authority must be able to bring the two together and decide whether the activity continues, changes, expands, or stops.

Risk functions often need business input to understand the activity, but the business cannot determine the risk conclusion simply because it owns the value case. Where disagreement remains material, the organization needs a way to record it and escalate the decision. How this is done will depend on organizational size, structure, and the consequence of the AI activity; RWR leaves the particular committee or reporting line to the organization.

7. Aster Vale Revisited: A Governance Decision

7.1 Original purpose and alternatives

Aster Vale did not pilot Navigator simply to make credit-memorandum preparation faster. The bank expected faster preparation to lead to faster decisions, lower applicant burden, and better access to individualized review for smaller businesses. Those expectations formed the original value case. Navigator was also not the bank's only option: it could retain the existing process, redesign the workflow without AI, use conventional automation, adopt a narrower AI-assisted tool, or continue with Navigator.

The value case also helped identify what the bank needed to learn. Saving preparation time would matter less if review and correction consumed the gain or if document preparation was not the main source of delay. Faster processing would matter differently if nonstandard applicants fared worse. And broader deployment would depend on whether source access, human verification, and pilot controls could survive scale.

7.2 What the bank learned

We now apply those distinctions to the pilot findings in Table 3. All figures and observations in the Aster Vale case are hypothetical. The table distinguishes task outcomes, broader value claims, and operating conditions that affect risk. Its final column explains how each observation bears on the next authorization decision.

Table 3. Aster Vale Navigator Pilot Findings

Dimension	Hypothetical pilot observation	Bearing on the decision
Financial	Preparation cost falls, but vendor, validation, quality-assurance, and integration costs absorb much of the saving.	Total net saving remains to be established.
Operational	Preparation time falls 28%; application-to-decision time falls 7%.	Task gains translate only partly to end-to-end gains.
Customer	Repeated requests fall 4%; abandonment changes from 13.0% to 12.7%.	The intended customer benefit remains weakly supported.
Access	No attributable approval effect; weak improvement for standardized small applications.	The access rationale remains unestablished.
Distribution	Standard digital applications improve; nonstandard-document processing time rises 6%.	An average gain conceals different effects across cases.
Workforce	Experienced users save time; verification burden persists; junior reliance increases.	User experience and review capacity affect value and risk.
Governance	Draft consistency improves; source accessibility falls from 94% to 64%.	Source access weakens despite a favorable drafting outcome.
Learning	Downstream bottlenecks, user differences, document limits, and scaling conditions become clearer.	The pilot identifies questions for the next decision.
Comparative	Non-AI redesign reduces preparation time 17% and decision time 6%.	Similar end-to-end gains require a fuller comparison.

The pilot supports a preparation-time gain and useful learning about Navigator’s operating conditions. Its broader implications are more limited: the end-to-end gain is smaller, the intended customer and access benefits remain unestablished, and effects differ across users and documents. The non-AI redesign further weakens a simple inference from preparation speed to a preference for Navigator. These observations change the value and risk questions the bank must answer before expansion.

7.3 The revised request

The bank is not simply asking to continue the pilot. Management now wants to expand Navigator across small-business lending, use it for larger loans and a broader range of documents, reduce quality-assurance coverage, and extend use to less-experienced underwriters.

The reason for expansion has changed as well. The original pilot emphasized faster decisions, lower applicant burden, and improved access. The revised case emphasizes processing capacity, avoided hiring, and the strategic value of developing AI capability.

At the same time, the proposed expansion increases exposure: more users, more heterogeneous documents, larger loans, and less oversight. The bank is therefore asking the pilot evidence to support a broader activity at the same time that both the value rationale and the risk conditions are changing.

7.4 Disposition

We begin with the requested expansion rather than with a general verdict on Navigator. The case supplies no finding that the restricted preparation use breaches an admissibility constraint, but the bank still has to confirm admissibility for the proposed scope. Risk permissibility is not yet established for expansion: source access has deteriorated, newer users rely more heavily on summaries, documents would become more varied, loan consequences would increase, and quality-assurance coverage would decline. Formal human decision authority does not resolve those dependencies.

The value case also requires revision. The original customer-and-access rationale received limited support, while the capacity-and-hiring rationale has not yet been tested against total implementation and control costs or the workflow alternative. The current record therefore supports withholding full expansion and investigating whether a narrower stage can be authorized. The bank could first restore source access, retain scope and quality-assurance limits, and specify the evidence needed for a direct comparison with workflow redesign. Any further live pilot must itself pass all three tests; learning potential alone cannot settle its risk permissibility.

If those conditions are satisfied, a controlled stage could test whether Navigator creates sufficient incremental value under workable controls. If they are not, restriction, suspension, or another feasible approach remains available. RWR's role is to make this sequence explicit: pilot success informs the next decision, while revised purpose and increased exposure require their own justification.

8. Discussion

8.1 What RWR adds

Aster Vale helps clarify the role of RWR. Existing risk management, benefits-realization, and AI-governance practices continue to do their respective work; RWR changes the decision they are asked to support. For the bank, staying within pilot risk limits would be only part of the evidence, as would the improvement in preparation time. Expansion requires the current value case and risk case to be reconsidered together in light of the changed rationale, changed exposure, and still-feasible alternatives.

That continuing authorization discipline is the main contribution of RWR. Competent decision makers may still reach different conclusions, but the approach identifies the questions to be answered, the evidence that should remain visible, and the circumstances that should reopen the decision. Its value therefore has to be tested in practice: does making these relationships explicit cause organizations to ask for different evidence, notice different dependencies, challenge changes in purpose earlier, or make different continuation and scaling decisions?

8.2 Why now

Aster Vale's problem becomes more likely when the activity and the evidence supporting it change on different schedules. Capabilities, vendors, users, costs, and alternatives may shift before the evidence for the prior decision matures. Authorization then needs an identifiable scope and conditions, together with observations that can trigger reconsideration. Rapid AI evolution strengthens the need for this discipline, although the same logic can serve other activities governed during incomplete learning.

8.3 Practical governance implications

RWR requires less institutional novelty than it may first appear. Existing governance can often accommodate the approach if someone owns and defends the value case, risk is assessed and challenged independently, and a decision maker has authority to change or withdraw authorization when the evidence no longer supports the current action. The harder requirement may be procedural rather than structural: both cases have to remain current enough to support a decision, material disagreement has to remain visible, and changes in purpose or evidence cannot disappear into routine monitoring.

Monitoring should follow the logic of the decision rather than place generic benefit and risk dashboards side by side. If faster document preparation is expected to create value through faster loan decisions, the bank should not stop at measuring preparation time; it should watch whether the downstream effect actually occurs. If source access is important to safe human review, deterioration in source access should matter to the risk case even if headline performance remains stable. The point is not to collect more indicators. It is to know which observations could change the next decision.

8.4 Limitations and research agenda

The practical benefit of RWR remains to be established in live organizational decisions. Our analysis identifies relationships to make explicit; empirical evaluation must determine whether doing so improves reasoning, changes action appropriately, or adds governance effort without sufficient benefit. Value justification will continue to involve stakeholder differences and judgment even with a better evidentiary record.

The separation between the value case and independent risk assessment may also be harder to maintain in practice than it appears on paper. Business sponsors often control information, resources, and organizational influence. The safeguards discussed earlier can support credible challenge, but they cannot guarantee it.

RWR also asks organizations to keep alternatives visible and to reconsider authorization after material change. That requires information, coordination, and management attention. For low-consequence or well-understood activities, the additional discipline may not be worth the burden.

The hypothetical Aster Vale case shows the reasoning sequence but supplies no evidence of effectiveness across organizations or sectors. That distinction guides the research questions below: first establish whether the approach changes decision-relevant reasoning, then examine whether those changes improve action and outcomes.

These limitations point directly to the empirical work ahead. Relevant questions may include:

- Risk discovery. Does purpose-directed inquiry identify different or earlier risks than category-based assessment alone?
- Decision effects. Does explicit separation of risk permissibility and value justification change continuation or scaling decisions?
- Persistence and burden. Does formal reauthorization after purpose or value revision prevent weak projects from persisting, or merely add procedural burden?
- Institutional design. What institutional arrangements produce credible value challenges without transferring business judgment to risk functions?
- Learning versus delay. Under what conditions does concurrent learning justify staged action rather than delay or termination?

A useful initial study could give comparable practitioner groups the same case evidence under ordinary governance and under RWR, with independent assessment of the requested evidence, identified dependencies, rationale, and disposition. Random assignment or counterbalancing would help separate the effect of RWR from additional reading, familiarity with the case, and order effects. Such an exercise could show whether the approach changes reasoning; evaluating whether it improves actual governance would require subsequent observation of decisions, outcomes, and costs.

9. Conclusion

Organizations increasingly deploy AI before they fully understand which outcomes an activity will produce, how much organizational value those outcomes will create, what resources and costs will be required to realize that value, or what risks will emerge. The governance challenge is therefore not only to assess risk and monitor benefits, but to decide whether evolving evidence about the value case and risk case continues to support the same organizational action.

We propose RWR as one way to organize this continuing decision. Independent risk identification provides the coverage floor, intended value adds purpose-directed emphasis, and the three authorization tests keep permissibility and value justification distinct. Alternatives remain visible, and material change can reopen the decision. The institutional shift is to treat authorization as a recurring decision rather than a permanent status. The practical question is simple: given what we know now, should this activity continue in its current form?

The resulting discipline keeps the original reason for accepting exposure open to review. Aster Vale may have learned enough to design a further controlled stage while still lacking justification for expansion. That distinction is the practical purpose of RWR: connect what the organization now knows about value and risk to what it proposes to do next. Whether the approach improves decisions and outcomes remains an empirical question.

AI-Use Disclosure

Generative AI tools—including ChatGPT and Codex, Claude, Gemini, Kimi, and Microsoft Copilot—assisted with literature discovery, drafting, editorial review, and document preparation. The author directed the research, reviewed and revised the resulting material, and takes full responsibility for the paper’s content.

References

- [1] Board of Governors of the Federal Reserve System, Office of the Comptroller of the Currency, and Federal Deposit Insurance Corporation. “Revised Guidance on Model Risk Management.” SR 26-2, 2026. <https://www.federalreserve.gov/supervisionreg/srletters/SR2602.htm>.
- [2] OECD, Boston Consulting Group, and INSEAD. *The Adoption of Artificial Intelligence in Firms: New Evidence for Policymaking*. OECD Publishing, 2025. <https://doi.org/10.1787/f9ef33c3-en>.
- [3] Committee of Sponsoring Organizations of the Treadway Commission (COSO). *Enterprise Risk Management— Integrating with Strategy and Performance*. 2017.
- [4] Office of the Comptroller of the Currency. “New, Modified, or Expanded Bank Products and Services: Risk Management Principles.” OCC Bulletin 2017-43, 2017. <https://www.occ.treas.gov/news-issuances/bulletins/2017/bulletin-2017-43.html>.
- [5] Federal Deposit Insurance Corporation, Board of Governors of the Federal Reserve System, and Office of the Comptroller of the Currency. “Interagency Guidance on Third-Party Relationships: Risk Management.” 2023. <https://www.fdic.gov/news/financial-institution-letters/2023/fil23029.html>.
- [6] Tabassi, Elham. *Artificial Intelligence Risk Management Framework (AI RMF 1.0)*. NIST AI 100-1, National Institute of Standards and Technology, 2023. <https://doi.org/10.6028/NIST.AI.100-1>.
- [7] Mäntymäki, Matti, Matti Minkkinen, Teemu Birkstedt, and Mika Viljanen. “Defining Organizational AI Governance.” *AI and Ethics* 2, 2022, 603–609. <https://doi.org/10.1007/s43681-022-00143-x>.
- [8] Birkstedt, Teemu, Matti Minkkinen, Anushree Tandon, and Matti Mäntymäki. “AI Governance: Themes, Knowledge Gaps and Future Agendas.” *Internet Research* 33(7), 2023, 133–167. <https://doi.org/10.1108/INTR-01-2022-0042>.
- [9] Mäntymäki, Matti, Matti Minkkinen, Markus Philipp Zimmer, Teemu Birkstedt, and Mika Viljanen. “Designing an AI Governance Framework: From Research-Based Premises to Meta-Requirements.” *ECIS 2023 Research Papers*, Paper 295, 2023. https://aisel.aisnet.org/ecis2023_rp/295/.
- [10] Bank of England and Financial Conduct Authority. *Artificial Intelligence in UK Financial Services—2024*. 2024. <https://www.bankofengland.co.uk/report/2024/artificial-intelligence-in-uk-financial-services-2024>.
- [11] Financial Stability Board. *The Financial Stability Implications of Artificial Intelligence*. 2024. <https://www.fsb.org/2024/11/the-financial-stability-implications-of-artificial-intelligence/>.
- [12] Infrastructure and Projects Authority. *Guide for Effective Benefits Management in Major Projects*. UK Cabinet Office, 2017. <https://www.gov.uk/government/publications/guide-for-effective-benefits-management-in-major-projects>.

- [13] Department of Finance, Northern Ireland. “Introduction to Benefits Management for Programmes and Projects.” <https://www.finance-ni.gov.uk/articles/introduction-benefits-management-programmes-and-projects>.
- [14] Mamabolo, Maisaka, and Carl Marnewick. “Benefits Realisation Management: Responsibilities and Challenges.” *South African Journal of Information Management* 24(1), 2022. <https://doi.org/10.4102/sajim.v24i1.1574>.
- [15] Williams, Byron K., Robert C. Szaro, and Carl D. Shapiro. *Adaptive Management: The U.S. Department of the Interior Technical Guide*. U.S. Department of the Interior, 2009. <https://www.usgs.gov/publications/adaptive-management-us-department-interior-technical-guide>.
- [16] Williams, Byron K. “Reducing Uncertainty About Objective Functions in Adaptive Management.” *Ecological Modelling* 225, 2012, 61–65. <https://doi.org/10.1016/j.ecolmodel.2011.11.009>.
- [17] Williams, Byron K., and Eleanor D. Brown. “Double Loop Learning in Adaptive Management: The Need, the Challenge, and the Opportunity.” *Environmental Management* 62, 2018, 995–1006. <https://doi.org/10.1007/s00267-018-1107-5>.
- [18] Haasnoot, Marjolijn, Jan H. Kwakkel, Warren E. Walker, and Judith ter Maat. “Dynamic Adaptive Policy Pathways: A Method for Crafting Robust Decisions for a Deeply Uncertain World.” *Global Environmental Change* 23(2), 2013, 485–498. <https://doi.org/10.1016/j.gloenvcha.2012.12.006>.
- [19] Ipsmiller, Edith, Keith D. Brouthers, and Desislava Dikova. “25 Years of Real Option Empirical Research in Management.” *European Management Review* 16(1), 2019, 55–68. <https://doi.org/10.1111/emre.12324>.
- [20] U.S. Food and Drug Administration. *Consideration of Uncertainty in Making Benefit-Risk Determinations in Medical Device Premarket Approvals, De Novo Classifications, and Humanitarian Device Exemptions*. 2019. <https://www.fda.gov/regulatory-information/search-fda-guidance-documents/consideration-uncertainty-making-benefit-risk-determinations-medical-device-premarket-approvals-de>.
- [21] Health and Safety Executive. “How HSE Regulates Risks to Health and Safety from Work Activities.” Updated 2023. <https://www.hse.gov.uk/enforce/how-hse-regulates/index.htm>.
- [22] Menzies, Tom, Gaelle Saint-Hilary, and Pavel Mozgunov. “A Comparison of Various Aggregation Functions in Multi-Criteria Decision Analysis for Drug Benefit-Risk Assessment.” *Statistical Methods in Medical Research* 31(5), 2022, 899–916. <https://doi.org/10.1177/09622802211072512>.
- [23] Stilgoe, Jack, Richard Owen, and Phil Macnaghten. “Developing a Framework for Responsible Innovation.” *Research Policy* 42(9), 2013, 1568–1580. <https://doi.org/10.1016/j.respol.2013.05.008>.
- [24] Jafari, Akbar Anbar, Cagri Ozcinar, and Gholamreza Anbarjafari. “Responsible AI: The Good, The Bad, The AI.” *arXiv preprint arXiv:2601.21095*, 2026. <https://arxiv.org/abs/2601.21095>.