

Return-Weighted Risk for Navigating an Evolving AI Landscape

Reconnecting Risk with Value in AI Risk Management

Honggao Cao
Independent Research
honggaoc@gmail.com
August 2026

Abstract

Organizations increasingly undertake AI activities while still learning how those activities create value, what resources they require, and what risks they introduce. Rapid changes in capabilities, configurations, user behavior, and feasible alternatives make point-in-time assessments short-lived, and initial business and risk cases often lose decision relevance as evidence evolves. Return-Weighted Risk (RWR) is an activity-level governance approach that keeps the evolving value case explicitly connected to an independently assessed risk case. Continuing authorization depends on three sequential tests: admissibility, risk permissibility, and value justification. Intended value guides additional risk inquiry without altering independent risk conclusions, and feasible alternatives remain relevant after initial approval. Material change in either case can reopen the decision. RWR reframes governance around a single question: whether current evidence about value and risk still supports the same organizational action. A hypothetical small-business-lending example illustrates the logic, and the paper identifies empirical questions for future validation.

Keywords: AI governance; AI risk management; return-weighted risk; activity-level governance; concurrent learning; continuing authorization; material change; value realization

1. The Governance Problem in an Evolving AI Environment

Organizations adopt artificial intelligence because they expect beneficial change—financial, operational, customer-related, workforce-related, strategic, scientific, public, or informational. The same activities can also create or amplify risks involving reliability, privacy, security, fairness, compliance, model behavior, third-party dependence, operational resilience, and human reliance. What makes the present environment especially difficult is that organizations must often learn about value and risk while the activity is already being developed, piloted, deployed, and changed.

This creates a moving governance target. AI capabilities, deployment configurations, user behavior, available alternatives, and organizational objectives can shift before the next decision point. A sound assessment today may therefore lose decision relevance sooner than governance processes built around more stable assumptions would expect.

When both sides are unsettled, the familiar intuition that risk can be characterized first and then used to condition a return decision becomes harder to apply. In many AI initiatives, organizations do not yet know enough about either the durable value pathway or the realized risk profile to treat one as a stable weight on the other. Waiting for complete knowledge can prevent useful learning; proceeding as though

initial business and risk cases were settled can allow authorization to persist after its justification has changed.

1.1 A motivating decision: Aster Vale

Consider Aster Vale Regional Bank, a hypothetical regional bank that completed a six-month pilot of Navigator, a vendor-developed generative-AI system used to organize documents and draft portions of credit memoranda for small-business loans. Navigator does not approve or decline applications, assign credit ratings, determine prices, or communicate directly with applicants. Underwriters retain formal decision authority and certify every completed memorandum.

The pilot appears successful from one perspective. Navigator reduces underwriter preparation time by 28 percent, experienced underwriters report spending less time sorting documents, and the system remains within its restricted scope. The bank identifies no autonomous credit decision, serious privacy incident, or explicit breach of approved pilot risk limits.

The broader evidence is mixed. Application-to-decision time declines only 7 percent because document preparation is not the principal source of delay. Repeated requests for applicant information fall only slightly, abandonment scarcely changes, and the anticipated improvement in credit access is not clearly established. A non-AI workflow redesign produces a similar improvement in end-to-end processing time.

Navigator performs best with standardized digital information; scanned, handwritten, multilingual, or unconventional records require more manual work. Some newer underwriters increasingly begin with Navigator summaries, and source accessibility deteriorates. The bank gains useful knowledge about generative AI while becoming more dependent on a vendor whose system it understands only partially.

Management nevertheless proposes expansion to all small-business lending teams and larger loans. Its revised rationale emphasizes processing capacity, avoided hiring, and strategic AI capability. The pilot, however, had originally been justified mainly by faster decisions, lower applicant burden, and improved access. The revised rationale may be legitimate, but it is materially different.

The expansion decision is therefore straightforward: given what the bank has learned about Navigator's value and risk, does the evidence still support broader deployment?

1.2 Return-weighted risk as a governance approach

The Aster Vale example raises the question at the center of this paper: How should organizations govern AI activities when they must learn about value realization and an evolving risk profile at the same time as they decide whether to continue, adapt, expand, restrict, suspend, or terminate those activities?

Our return-weighted risk (RWR) approach is intended to address that question. Instead of asking only whether value and risk are being managed, we ask whether current evidence about the two continues to support the same organizational action. RWR changes the decision in several ways. Intended value helps direct additional risk inquiry without determining the risk conclusion. Acceptable risk alone does not justify continuation. Feasible alternatives remain relevant after initial approval. And material change can reopen the authorization decision.

Section 2 explains why rapid AI evolution makes point-in-time assessments unusually transient. Section 3 positions the problem against existing governance foundations. Sections 4 through 6 introduce RWR and develop its authorization and learning logic. Section 7 revisits Aster Vale. Section 8 discusses implications, limitations, and research questions. Section 9 concludes.

2. Why AI Makes the Problem More Acute

The difficulty described in Section 1 is intensified by the way AI activities change while organizations are still learning about them. Assessments can become stale, the activity itself can evolve, favorable outcomes may not translate into organizational value, and risk can shift across users and operating conditions. Evidence that supported one decision may therefore be less useful for the next.

2.1 Rapid change makes point-in-time assessment transient

AI changes not only what organizations know, but how long that knowledge remains decision-relevant. A point-in-time assessment can be carefully performed and still become stale as capabilities expand, vendors update systems, deployment patterns shift, costs change, or feasible alternatives improve. In some cases, the activity being assessed is no longer quite the same activity by the time the next decision is made.

The problem affects both sides of the decision. Expected value can change as capabilities and alternatives shift, while new evidence can change what the organization believes has actually been realized. Risk can also change as configurations, users, and operating conditions create different exposures. Continuing governance therefore cannot assume that a sound initial assessment remains stable simply because it was sound when made.

In April 2026, U.S. banking agencies explicitly excluded generative and agentic AI from revised model-risk-management guidance because those technologies were considered “novel and rapidly evolving” ([1]). The agencies also announced plans for further information gathering focused on banks’ use of AI, including generative and agentic systems.

Established governance principles still apply, and banks remain responsible for appropriate internal governance and controls. The episode nevertheless illustrates the difficulty of assuming that a conventional point-in-time model-risk framework can fully settle governance for a rapidly changing AI activity.

2.2 Purpose and the deployed activity can change

A formal change in scope is usually easy to see: an AI system may move from document preparation to recommendation, from a restricted user group to an enterprise workflow, or from an internal aid to customer-facing interaction. More difficult is the gradual change that occurs when users begin relying on the system differently, prompts or retrieval change, the vendor updates the product, or an internal aid begins shaping decisions more heavily than originally intended. Governance labels can remain stable while the effective use changes underneath them.

This can happen because the deployed activity is more than the model.

Data, retrieval, prompts, tools, human review, workflow, vendor changes, controls, and operating context all influence what happens in use. Aster Vale, for example, is not governing Navigator in isolation from its underwriters, source access, document types, verification practices, and workflow. Evidence from a pilot therefore describes a configured activity under particular conditions; it does not automatically establish how the same underlying model will perform when those conditions change.

2.3 From outcome to value

An observed outcome is not automatically value. Navigator can reduce document-preparation time without producing a comparable improvement in end-to-end lending performance. More generally, a task-level improvement matters only to the extent that it advances a relevant organizational objective or creates value for an affected stakeholder.

Even when organizational outcomes improve, attribution remains difficult. The improvement may reflect AI use, training, workflow redesign, user selection, policy changes, or interactions among them. Evidence from the Organisation for Economic Co-operation and Development (OECD) that many firms struggle to estimate AI return on investment is consistent with this difficulty ([2]). Organizations should therefore be cautious about treating an observed improvement as realized AI value before both the downstream consequence and the basis for attribution are reasonably supported.

2.4 Risk can shift across users and conditions

The same distributed structure complicates the risk side. An activity can perform well on average while creating different exposure for different users or cases. Experienced employees may save time while newer employees become more reliant on AI summaries; standardized documents may work well while nonstandard records require more verification or produce poorer results. The relevant risk therefore depends not only on model behavior, but also on who uses the activity, on what inputs, under what controls, and at what scale. Distribution belongs in both the value and risk cases.

2.5 Evidence may not travel

A successful pilot is evidence about the pilot, not automatically about the proposed expansion. Scaling changes users, volumes, oversight, data, incentives, consequences, and sometimes the activity itself. Evidence from experienced underwriters working with standardized documents under close pilot controls, for example, may not support broader use by less-experienced users on more varied documents with reduced quality assurance. The more the proposed conditions differ from those under which the evidence was generated, the more carefully the organization should ask whether that evidence still applies. This portability problem affects value and risk at the same time.

2.6 Governing while learning

Taken together, these conditions create a practical problem: organizations may have to decide how to proceed before either the value case or the risk picture is fully settled. Waiting for complete knowledge

may prevent useful learning; proceeding as though the initial assessment remains valid can create a different problem.

The situation resembles crossing a river by testing the footing ahead. The destination may be understood even when the complete path is not, and each controlled step produces information for the next. The analogy does not justify blind experimentation. Each step still has to be limited enough that the organization can learn without taking exposure it would not otherwise permit. RWR adopts this logic: move carefully, learn from what happens, and be prepared to decide again.

Organizations therefore need a way to move carefully, learn from experience, and revisit the decision as the evidence changes.

Existing approaches already provide much of the thinking needed to connect risk, value, and learning.

3. Foundations for the RWR Approach

The problem described above is not entirely new, and organizations do not have to solve it from first principles. Existing approaches already tell us a great deal about how to govern risk, pursue value, and make decisions while learning. RWR builds on those foundations. The purpose of this section is to show the ideas on which the approach relies and then narrow the discussion to RWR's particular focus.

3.1 Risk, value, and organizational purpose

Risk management has long been connected to organizational objectives rather than treated as an isolated control function. The Committee of Sponsoring Organizations of the Treadway Commission (COSO) enterprise risk management framework integrates risk with strategy and performance ([3]). In banking, Office of the Comptroller of the Currency (OCC) new-activity guidance explicitly connects strategy, risks and returns, alternatives, performance criteria, exit strategies, and ongoing review ([4]). Revised model-risk and third-party guidance adds expectations for monitoring, validation, accountability, and continuing fitness for purpose ([1], [5]).

AI governance follows much the same direction. The National Institute of Standards and Technology (NIST) Artificial Intelligence Risk Management Framework (AI RMF) asks organizations to consider intended purpose, beneficial uses, assumptions, limitations, and positive and negative impacts while managing risk across the lifecycle ([6]). Organizational AI-governance research similarly connects AI use with organizational strategy, objectives, values, and obligations ([7]).

The field is still developing. Reviews identify fragmentation and implementation gaps in organizational AI governance, while design-oriented work has proposed more practice-oriented requirements ([8], [9]). Financial-sector surveys report substantial expectations for AI alongside concerns about data, privacy, security, bias, model complexity, third-party dependence, and skills ([10], [11]). OECD evidence adds a practical difficulty: firms often struggle to estimate AI return on investment, and experimentation does not guarantee success ([2]). Taken together, this work shows that risk can be considered alongside the organizational purpose for undertaking an activity, even when the expected value remains difficult to establish.

3.2 Value realization

Benefits realization management (BRM) addresses what happens after a business case is approved. Expected benefits should be specified, assigned to owners, measured against baselines, followed beyond project delivery, and reconciled with the business case ([12], [13]). Practice is not always easy; ownership, resources, and accountability can break down ([14]). Approval of an expected value case is therefore not evidence that the value was later realized. Aster Vale illustrates the point: Navigator produced a clear task-level improvement while the broader customer and access value remained uncertain.

3.3 Learning under uncertainty

Benefits realization does not eliminate the problem that some uncertainty can only be reduced through experience. Adaptive management addresses that problem more directly: stakeholders, objectives, alternatives, models, monitoring, and observed responses are used to update decisions as evidence emerges ([15]). Later work recognizes that learning can change not only technical understanding but also objectives or institutional decision assumptions ([16], [17]). Dynamic Adaptive Policy Pathways and real-options reasoning add related ideas of staged commitment, adaptation triggers, and options to expand or abandon as information improves ([18], [19]).

Incomplete knowledge, then, does not always require either full commitment or complete delay. Limited action can sometimes be used to learn, provided that the exposure required to learn remains acceptable.

3.4 Constraints on value-risk trade-offs

Expected value cannot justify every level or type of risk. Existing benefit-risk approaches already recognize this limit. Medical-device regulation can consider uncertainty while maintaining safety and effectiveness requirements ([20]). The Health and Safety Executive (HSE) uses a tolerability-of-risk framework alongside a continuing duty to reduce risk so far as is reasonably practicable ([21]). Multicriteria benefit-risk research warns that simple aggregation can produce counterintuitive recommendations ([22]). Responsible-innovation scholarship adds anticipation, reflexivity, inclusion, and responsiveness, extending attention beyond organizational return alone ([23]). Together, these traditions support considering constraints and risk limits before asking whether the remaining exposure is justified by value.

3.5 The Focus of RWR

Taken together, these traditions supply much of the thinking RWR needs. Risk governance connects exposure to organizational objectives and limits. Benefits realization keeps attention on whether expected value actually materializes. Adaptive approaches allow decisions to evolve as evidence accumulates. Benefit-risk traditions place limits on what expected value can justify.

Recent AI-governance work also brings value creation and risk mitigation into the same conversation. The recent Paradox-based Responsible AI Governance (PRAIG) preprint is especially close to RWR in treating that relationship as a persistent dynamic tension rather than a static trade-off ([24]). PRAIG

addresses organizational or portfolio posture toward those tensions; RWR takes a narrower activity-level focus.

Consider Aster Vale again. The pilot produced favorable outcomes, but the original value rationale is only partly supported. Management now proposes a different rationale and broader exposure, while a non-AI alternative remains feasible. The question for RWR is whether, taken together, the current value case and independently assessed risk case still support the same action.

That continuing authorization decision is the focus of RWR. Table 1 summarizes how this focus relates to several adjacent governance perspectives.

Table 1. Governance perspectives on the Aster Vale expansion decision

Perspective	Primary question in the Aster Vale decision
Risk management	Is residual risk acceptable for the proposed scope and controls?
Benefits realization	Are intended benefits defined, owned, measured, and being realized?
Adaptive management	What should be learned next, and how should action adapt as evidence changes?
PRAIG	What organizational posture should manage the persistent tension between AI value creation and risk mitigation?
RWR	Do the current value case and current risk case, relative to feasible alternatives, support the same proposed action?

The perspectives are complementary. OCC guidance, for example, could support a similar ultimate disposition ([4]). RWR focuses specifically on the point at which the evolving value case and risk case must support the same proposed action.

4. Return-Weighted Risk: The Governance Approach

4.1 RWR in plain language

In practice, RWR starts from a simple situation. An organization approved an AI activity for a reason. As the activity operates, it learns more about what the activity actually delivers and what risks it actually creates. RWR asks whether the value story still holds and whether the risk picture has changed. More formally, it keeps the evolving value case connected to an independently assessed risk case and asks whether both still support the same action. If either case changes materially, the organization decides again rather than merely continuing to monitor.

4.2 Why “Return-Weighted Risk”?

In conventional investment reasoning, risk helps condition the return decision. Once expected return and risk are sufficiently understood, the question is whether the expected return justifies the risk being taken. RWR starts from a different situation. For many evolving AI activities, neither side is sufficiently settled: the organization is still learning what value the activity can produce, what resources and costs are required to realize that value, and what risks will emerge as the activity changes.

RWR therefore turns the familiar phrase around. Intended value helps tell the organization where additional risk inquiry may be needed, while independent risk assessment prevents the value case from determining the risk conclusion. The reversal is conceptual, not mathematical; RWR is not a formula for converting value and risk into a common score. We use return in the name to preserve the connection with risk-weighted return and to represent the economic side of the decision. In practice, RWR works through a value case because relevant value may be financial, operational, customer-related, workforce-related, strategic, or otherwise difficult to reduce to a single return measure.

4.3 The value case

Every AI activity is undertaken for a reason. The organization expects something useful to happen: work may become faster, service may improve, costs may fall, capacity may increase, customers may benefit, or a strategic capability may be developed. The value case makes that reason explicit and keeps it open to challenge as evidence accumulates. It preserves the distinction developed in Section 2.3 between an observed outcome and the value that outcome creates for the organization or an affected stakeholder.

The value case therefore keeps enough information to understand what the organization is trying to achieve, for whom, relative to what alternative, at what cost, and with what evidence. Relevant details will vary with the decision, but usually include the important outcomes, stakeholders, resources and costs, alternatives, and supporting evidence. What the organization intends to create is not necessarily what it can reasonably expect, and neither is the same as what it later observes or can support as realized value. As evidence changes, expectations should change with it. Attribution requires similar care: improvement under an AI-enabled activity does not by itself establish how much of that improvement was caused by AI rather than workflow, training, user selection, or other complementary changes.

4.4 The risk case

The value case explains why the organization wants to undertake the activity. The risk case asks what can go wrong, who may be affected, how serious the exposure may be, and whether controls make the remaining risk acceptable. RWR uses risk in the narrower sense of possible adverse effects under uncertainty.

Cost should not be hidden inside the risk case. Expected spending on technology, people, verification, controls, or other resources belongs in the economic evaluation of the activity. Uncertainty about those costs may create risk, but the expected cost itself is not risk. The risk case also remains independent of the value case. Knowing why an activity matters can help identify dependencies or consequences that deserve additional purpose-directed inquiry, but expected value cannot determine the risk conclusion.

4.5 Coupled governance

The two cases remain distinct but meet at one authorization decision. Business remains responsible for explaining why the activity should be undertaken and for supporting the value case. Risk functions independently assess and challenge the risk; neither substitutes for the other. The authorization body then asks whether both cases support the proposed action. An existing governance structure can perform this role when it has the necessary authority and challenge capability.

5. Continuing Authorization Under RWR

Under RWR, authorization does not end with the initial approval. As the value and risk cases change, the organization must be prepared to ask again whether the activity should continue in its current form. We organize that decision around three questions.

5.1 Three tests

RWR uses three sequential tests for the authorization decision. Their formal questions appear below, and Table 2 summarizes the decision logic.

- **Admissibility:** Does the action satisfy applicable legal, ethical, policy, and non-compensable constraints?
- **Risk permissibility:** Does independently assessed residual risk remain within authorized limits for the proposed scope and conditions?
- **Value justification:** Does the current value case justify undertaking that permissible residual risk?

The order matters. Some actions should not proceed regardless of expected value; they fail the admissibility test. Other actions may be admissible but still carry more residual risk than the organization is prepared to accept. Only after an action passes those two tests does RWR ask whether the value case is strong enough to justify taking the permissible risk. Passing the risk test therefore does not mean approve; it means the activity is eligible for the value-justification decision.

Table 2. RWR three-test authorization sequence

Test	Plain-language question	If not satisfied
Admissibility	May this action proceed at all?	Do not authorize.
Risk permissibility	Is residual risk within permitted limits for the proposed conditions?	Do not authorize in the proposed form.
Value justification	Does the current value case justify undertaking that permissible risk?	Do not authorize; adapt, narrow, or reconsider.

5.2 Value justification as governed judgment

Value justification is not always a calculation. Some AI activities have effects that can be expressed in money; others involve customer experience, workforce effects, access, strategic capability, or learning that may not be meaningfully reduced to the same unit. RWR does not solve that problem by inventing a universal score. The authorization body must make a judgment from the evidence available—what value is expected or has been realized, for whom, at what cost, relative to what alternatives, and with what uncertainty. Calling this a governed judgment does not mean using managerial intuition alone. The decision should be made by the appropriate authority, on an explicit evidentiary record, with challenge and accountability.

The strength of evidence required also depends on the decision being made. Evidence that is good enough to continue a restricted pilot may not be good enough to expand the same activity across the enterprise. The point is not to label evidence mechanically as sufficient or insufficient, but to ask whether it is sufficient for this decision.

5.3 Avoiding unnecessary exposure

Passing all three tests may leave more than one feasible option. Several approaches may be admissible, risk-permissible, and supported by a credible value case. Rather than reduce every form of value and risk to a common score and solve for an optimum, RWR asks a simpler question: is there a feasible way to pursue the justified purpose without taking unnecessary residual risk? For Aster Vale, that means asking why broader Navigator exposure is necessary if workflow redesign can deliver much of the same operational improvement with less AI-related exposure.

5.4 Alternatives remain live

The comparison does not end when the initial business case is approved. Alternatives change too. A non-AI workflow may improve, another AI option may become available, costs may shift, or experience may show that the original bottleneck was somewhere else. RWR therefore keeps feasible alternatives in view. An activity should not continue merely because it produces some value if another feasible approach can pursue much the same purpose at lower cost, with stronger evidence, or with less exposure.

The comparison need not rebuild the full business and risk case for every alternative at every review. It should be refreshed when decision-relevant evidence changes materially. Navigator's preparation-time improvement, for example, looks different once Aster Vale learns that workflow redesign produces almost the same end-to-end improvement.

5.5 Material change and reauthorization

Not every change requires reauthorization. For RWR, a change becomes material when it could plausibly change the outcome of one of the three tests. A change in purpose may do this. So may a substantial change in users, scale, operating conditions, available alternatives, evidence, or risk exposure. The trigger does not have to take one form: some changes can be tied to quantitative thresholds, others require judgment, and high-uncertainty activities may warrant scheduled reconsideration even when no single event crosses a threshold. Reauthorization does not imply that the revised activity is improper. It recognizes that the organization may now be accepting risk for a different reason or under materially different conditions.

For example, if a document-preparation AI is re-scoped to recommend credit decisions, the activity has changed materially even if the underlying vendor or model has not.

6. Governing Concurrent Learning

6.1 The operating cycle

Sometimes the three tests do not produce a simple continue-or-stop decision because important uncertainty remains. The organization may not know enough to justify broader deployment, but it may be able to learn what it needs through a smaller, controlled step. RWR does not treat uncertainty itself as a reason to experiment: the action required to learn must still be admissible and risk-permissible, and the

expected value of the learning must justify taking that exposure. When those conditions are met, authorization can become a learning cycle rather than a one-time decision. Figure 1 summarizes that cycle.




Figure 1. RWR operating cycle

The organization states what it is trying to achieve, identifies what it still needs to know, acts within approved limits, observes what happens, updates the value and risk cases, and decides again. The cycle does not override the three tests; a learning stage is itself an authorized action. Nor does the organization have to wait for the planned stage to end. If something changes materially, the authorization decision can be reopened immediately.

6.2 What needs to be learned

Learning has value in RWR only when it can inform a decision. Before beginning a learning stage, the organization should be clear about what it needs to know and what it would do differently depending on the answer. That usually means identifying the question being tested, what will be observed, what comparison matters, and what result would change the next authorization decision.

Aster Vale, for example, does not need another pilot simply to accumulate more experience with Navigator. It needs evidence that bears on specific decisions: whether source access can be restored, whether less-experienced users can use the system without excessive reliance, and whether Navigator produces enough incremental value over workflow redesign to justify broader exposure. More information is useful only when it can change what the organization does.

6.3 Separate evidence, coupled decisions

Value and risk may be learned from some of the same experience, but they remain different questions. Strong control evidence does not prove value, and strong value evidence does not repair a weak risk case. The same data can contribute to both cases. What matters is that the organization remains clear about what the evidence supports, who interprets it, and how it affects the decision.

The cases can remain distinct while informing one another. Evidence about value may reveal a dependency that deserves additional risk inquiry. Emerging risk may make the expected value harder or more costly to realize. One case should not substitute for the other. Neither stream determines the other; both affect authorization.

6.4 Decision authority and challenge

RWR can operate within an existing organizational structure. The value case needs an owner, the risk case needs credible independent assessment and challenge, and someone with appropriate authority must be able to bring the two together and decide whether the activity continues, changes, expands, or stops.

Risk functions often need business input to understand the activity, but the business cannot determine the risk conclusion simply because it owns the value case. Where disagreement remains material, the organization needs a way to record it and escalate the decision. How this is done will depend on organizational size, structure, and the consequence of the AI activity; RWR leaves the particular committee or reporting line to the organization.

7. Aster Vale Revisited: A Governance Decision

7.1 Original purpose and alternatives

Aster Vale did not pilot Navigator simply to make credit-memorandum preparation faster. The bank expected faster preparation to lead to faster decisions, lower applicant burden, and better access to individualized review for smaller businesses. Those expectations formed the original value case. Navigator was also not the bank's only option: it could retain the existing process, redesign the workflow without AI, use conventional automation, adopt a narrower AI-assisted tool, or continue with Navigator.

The value case also helped identify what the bank needed to learn. Saving preparation time would matter less if review and correction consumed the gain or if document preparation was not the main source of delay. Faster processing would matter differently if nonstandard applicants fared worse. And broader deployment would depend on whether source access, human verification, and pilot controls could survive scale.

7.2 What the bank learned

Table 3 summarizes what Aster Vale learned during the restricted Navigator pilot. The assessment labels are qualitative reading aids rather than points on a common scale: some describe evidence strength (for example, directional or insufficient), while others describe the pattern of observed effects (for example, mixed or uneven).

Table 3. Aster Vale Navigator Pilot Findings

Dimension	Pilot observation	Assessment
Financial	Preparation cost falls, but vendor, validation, QA, and integration absorb much of the saving.	Directional
Operational	Preparation time falls 28%; application-to-decision time falls 7%.	Mixed
Customer	Repeated requests fall 4%; abandonment changes from 13.0% to 12.7%.	Insufficient
Access	No attributable approval effect; weak improvement for standardized small applications.	Insufficient
Distribution	Standard digital applications improve; nonstandard-document processing time rises 6%.	Uneven
Workforce	Experienced users save time; verification burden persists; junior reliance increases.	Mixed
Governance	Draft consistency improves; source accessibility falls from 94% to 64%.	Mixed / weakened
Learning	Downstream bottlenecks, user differences, document limits, and scaling conditions become clearer.	Supported for pilot
Comparative	Non-AI redesign reduces preparation time 17% and decision time 6%.	AI advantage limited

The pilot worked, but not quite in the way the original value case anticipated. Navigator clearly reduced preparation time, and the pilot taught the bank useful things about where the system works and where it does not. But faster preparation translated into only a small improvement in end-to-end decision time. The expected customer and access benefits were not established, some effects were uneven, and a non-AI workflow redesign produced nearly the same improvement in overall processing time. The pilot was therefore neither a failure nor an obvious case for expansion. It changed what the bank knew.

7.3 The revised request

The bank is not simply asking to continue the pilot. Management now wants to expand Navigator across small-business lending, use it for larger loans and a broader range of documents, reduce quality-assurance coverage, and extend use to less-experienced underwriters.

The reason for expansion has changed as well. The original pilot emphasized faster decisions, lower applicant burden, and improved access. The revised case emphasizes processing capacity, avoided hiring, and the strategic value of developing AI capability.

At the same time, the proposed expansion increases exposure: more users, more heterogeneous documents, larger loans, and less oversight. The bank is therefore asking the pilot evidence to support a broader activity at the same time that both the value rationale and the risk conditions are changing.

7.4 Disposition

RWR does not point to either automatic expansion or abandonment. Nothing in the case suggests that the restricted preparation use has become inadmissible. The proposed expansion is different. Source access has deteriorated, less-experienced users rely more heavily on Navigator, the document population would broaden, loan consequences would increase, and management proposes less quality-assurance coverage. The pilot therefore does not yet establish that residual risk remains permissible at the proposed scale.

The value case is also unsettled. The original customer-and-access rationale received limited support, while the new efficiency-centered rationale has not yet been tested against total implementation and control costs or against the non-AI workflow alternative. RWR would therefore support another controlled stage rather than full expansion. The bank could restore source access, retain scope limits, compare Navigator directly with workflow redesign, examine workforce and distributional effects, and specify what evidence would justify broader deployment.

The important point is not that RWR produces a uniquely correct answer. It changes what must be reconsidered before the bank treats a successful pilot as justification for broader use.

8. Discussion

8.1 What RWR adds

Aster Vale helps clarify the role of RWR. Existing risk management, benefits-realization, and AI-governance practices continue to do their respective work; RWR changes the decision they are asked to support. For the bank, staying within pilot risk limits would be only part of the evidence, as would the improvement in preparation time. Expansion requires the current value case and risk case to be reconsidered together in light of the changed rationale, changed exposure, and still-feasible alternatives.

That continuing authorization discipline is the main contribution of RWR. Competent decision makers may still reach different conclusions, but the approach identifies the questions to be answered, the evidence that should remain visible, and the circumstances that should reopen the decision. Its value therefore has to be tested in practice: does making these relationships explicit cause organizations to ask for different evidence, notice different dependencies, challenge changes in purpose earlier, or make different continuation and scaling decisions?

8.2 Why now

Uncertainty by itself is not the reason for RWR; organizations have always made decisions under uncertainty. What makes the present AI environment particularly difficult is the combination of rapid change and incomplete learning. Capabilities change, configurations change, users adapt, vendors update systems, alternatives improve, and organizational purposes shift—often before the evidence supporting the previous decision has matured. In a more stable setting, it is easier to imagine characterizing risk and using it to condition a return decision. In an evolving AI activity, the organization may still be learning about both sides at once. RWR is intended for that situation. Established governance approaches remain

relevant and provide much of its foundation; the changing environment simply puts greater weight on revisiting authorization as the evidence changes.

8.3 Practical governance implications

RWR requires less institutional novelty than it may first appear. Existing governance can often accommodate the approach if someone owns and defends the value case, risk is assessed and challenged independently, and a decision maker has authority to change or withdraw authorization when the evidence no longer supports the current action. The harder requirement may be procedural rather than structural: both cases have to remain current enough to support a decision, material disagreement has to remain visible, and changes in purpose or evidence cannot disappear into routine monitoring.

Monitoring should follow the logic of the decision rather than place generic benefit and risk dashboards side by side. If faster document preparation is expected to create value through faster loan decisions, the bank should not stop at measuring preparation time; it should watch whether the downstream effect actually occurs. If source access is important to safe human review, deterioration in source access should matter to the risk case even if headline performance remains stable. The point is not to collect more indicators. It is to know which observations could change the next decision.

8.4 Limitations and research agenda

The most important limitation is also the most important next question: RWR has not yet been tested in live organizational decisions. We do not know whether making the value-risk relationship explicit will improve decisions or simply add another layer of governance effort. Difficult value judgments will remain. Different stakeholders may value outcomes differently, some effects cannot be reduced to a common measure, and reasonable decision makers may disagree about whether the value case is strong enough to justify permissible risk. RWR makes that judgment more visible and structured, not objective.

The separation between the value case and independent risk assessment may also be harder to maintain in practice than it appears on paper. Business sponsors often control information, resources, and organizational influence. The safeguards discussed earlier can support credible challenge, but they cannot guarantee it.

RWR also asks organizations to keep alternatives visible and to reconsider authorization after material change. That requires information, coordination, and management attention. For low-consequence or well-understood activities, the additional discipline may not be worth the burden.

And Aster Vale remains an illustration. It shows how the logic can work, not whether it works equally well across organizations, sectors, or higher-consequence AI uses.

These limitations point directly to the empirical work ahead. Relevant questions may include:

- **Risk discovery.** Does purpose-directed inquiry identify different or earlier risks than category-based assessment alone?
- **Decision effects.** Does explicit separation of risk permissibility and value justification change continuation or scaling decisions?

- **Persistence and burden.** Does formal reauthorization after purpose or value revision prevent weak projects from persisting, or merely add procedural burden?
- **Institutional design.** What institutional arrangements produce credible value challenges without transferring business judgment to risk functions?
- **Learning versus delay.** Under what conditions does concurrent learning justify staged action rather than delay or termination?

A useful first test does not require a large field study. Practitioners could assess the same AI-use case first under their ordinary governance practice and then using the RWR three-test approach. Differences in requested evidence, identified dependencies, rationale, and disposition could then be compared. Such an exercise would not show that RWR improves real-world outcomes, but it would tell us whether the approach changes decision-relevant reasoning at all. If it does not, there may be little reason to proceed to more costly field studies.

9. Conclusion

Organizations increasingly deploy AI before they fully understand which outcomes an activity will produce, how much organizational value those outcomes will create, what resources and costs will be required to realize that value, or what risks will emerge. The governance challenge is therefore not only to assess risk and monitor benefits, but to decide whether evolving evidence about the value case and risk case continues to support the same organizational action.

We propose RWR as one way to organize this continuing decision. Independent risk identification provides the coverage floor, intended value adds purpose-directed emphasis, and the three authorization tests keep permissibility and value justification distinct. Alternatives remain visible, and material change can reopen the decision. The institutional shift is to treat authorization as a recurring decision rather than a permanent status. The practical question is simple: given what we know now, should this activity continue in its current form?

This paper defines the RWR approach and the problem it is intended to address. Whether RWR improves governance in practice remains an empirical question.

AI-use disclosure. Generative AI tools—including ChatGPT and Codex, Claude, Gemini, Kimi, and Microsoft Copilot—assisted with literature discovery, drafting, editorial review, and document preparation. The author directed the research, reviewed and revised the resulting material, and takes full responsibility for the paper’s content.

References

- [1] Board of Governors of the Federal Reserve System, Office of the Comptroller of the Currency, and Federal Deposit Insurance Corporation. “Revised Guidance on Model Risk Management.” SR 26-2, 2026. <https://www.federalreserve.gov/supervisionreg/srletters/SR2602.htm>.
- [2] OECD, Boston Consulting Group, and INSEAD. The Adoption of Artificial Intelligence in Firms: New Evidence for Policymaking. OECD Publishing, 2025. <https://doi.org/10.1787/f9ef33c3-en>.

- [3] Committee of Sponsoring Organizations of the Treadway Commission (COSO). Enterprise Risk Management—Integrating with Strategy and Performance. 2017.
- [4] Office of the Comptroller of the Currency. “New, Modified, or Expanded Bank Products and Services: Risk Management Principles.” OCC Bulletin 2017-43, 2017. <https://www.occ.treas.gov/news-issuances/bulletins/2017/bulletin-2017-43.html>.
- [5] Federal Deposit Insurance Corporation, Board of Governors of the Federal Reserve System, and Office of the Comptroller of the Currency. “Interagency Guidance on Third-Party Relationships: Risk Management.” 2023. <https://www.fdic.gov/news/financial-institution-letters/2023/fil23029.html>.
- [6] Tabassi, Elham. Artificial Intelligence Risk Management Framework (AI RMF 1.0). NIST AI 100-1, National Institute of Standards and Technology, 2023. <https://doi.org/10.6028/NIST.AI.100-1>.
- [7] Mäntymäki, Matti, Matti Minkkinen, Teemu Birkstedt, and Mika Viljanen. “Defining Organizational AI Governance.” *AI and Ethics* 2, 2022, 603–609. <https://doi.org/10.1007/s43681-022-00143-x>.
- [8] Birkstedt, Teemu, Matti Minkkinen, Anushree Tandon, and Matti Mäntymäki. “AI Governance: Themes, Knowledge Gaps and Future Agendas.” *Internet Research* 33(7), 2023, 133–167. <https://doi.org/10.1108/INTR-01-2022-0042>.
- [9] Mäntymäki, Matti, Matti Minkkinen, Markus Philipp Zimmer, Teemu Birkstedt, and Mika Viljanen. “Designing an AI Governance Framework: From Research-Based Premises to Meta-Requirements.” *ECIS 2023 Research Papers*, Paper 295, 2023. https://aisel.aisnet.org/ecis2023_rp/295/.
- [10] Bank of England and Financial Conduct Authority. Artificial Intelligence in UK Financial Services—2024. 2024. <https://www.bankofengland.co.uk/report/2024/artificial-intelligence-in-uk-financial-services-2024>.
- [11] Financial Stability Board. The Financial Stability Implications of Artificial Intelligence. 2024. <https://www.fsb.org/2024/11/the-financial-stability-implications-of-artificial-intelligence/>.
- [12] Infrastructure and Projects Authority. Guide for Effective Benefits Management in Major Projects. UK Cabinet Office, 2017. <https://www.gov.uk/government/publications/guide-for-effective-benefits-management-in-major-projects>.
- [13] Department of Finance, Northern Ireland. “Introduction to Benefits Management for Programmes and Projects.” <https://www.finance-ni.gov.uk/articles/introduction-benefits-management-programmes-and-projects>.
- [14] Mamabolo, Maisaka, and Carl Marnewick. “Benefits Realisation Management: Responsibilities and Challenges.” *South African Journal of Information Management* 24(1), 2022. <https://doi.org/10.4102/sajim.v24i1.1574>.
- [15] Williams, Byron K., Robert C. Szaro, and Carl D. Shapiro. Adaptive Management: The U.S. Department of the Interior Technical Guide. U.S. Department of the Interior, 2009. <https://www.usgs.gov/publications/adaptive-management-us-department-interior-technical-guide>.
- [16] Williams, Byron K. “Reducing Uncertainty About Objective Functions in Adaptive Management.” *Ecological Modelling* 225, 2012, 61–65. <https://doi.org/10.1016/j.ecolmodel.2011.11.009>.
- [17] Williams, Byron K., and Eleanor D. Brown. “Double Loop Learning in Adaptive Management: The Need, the Challenge, and the Opportunity.” *Environmental Management* 62, 2018, 995–1006. <https://doi.org/10.1007/s00267-018-1107-5>.
- [18] Haasnoot, Marjolijn, Jan H. Kwakkel, Warren E. Walker, and Judith ter Maat. “Dynamic Adaptive Policy Pathways: A Method for Crafting Robust Decisions for a Deeply Uncertain World.” *Global Environmental Change* 23(2), 2013, 485–498. <https://doi.org/10.1016/j.gloenvcha.2012.12.006>.
- [19] Ipsmiller, Edith, Keith D. Brouthers, and Desislava Dikova. “25 Years of Real Option Empirical Research in Management.” *European Management Review* 16(1), 2019, 55–68. <https://doi.org/10.1111/emre.12324>.

- [20] U.S. Food and Drug Administration. Consideration of Uncertainty in Making Benefit-Risk Determinations in Medical Device Premarket Approvals, De Novo Classifications, and Humanitarian Device Exemptions. 2019. <https://www.fda.gov/regulatory-information/search-fda-guidance-documents/consideration-uncertainty-making-benefit-risk-determinations-medical-device-premarket-approvals-de>.
- [21] Health and Safety Executive. “How HSE Regulates Risks to Health and Safety from Work Activities.” Updated 2023. <https://www.hse.gov.uk/enforce/how-hse-regulates/index.htm>.
- [22] Menzies, Tom, Gaelle Saint-Hilary, and Pavel Mozgunov. “A Comparison of Various Aggregation Functions in Multi-Criteria Decision Analysis for Drug Benefit-Risk Assessment.” *Statistical Methods in Medical Research* 31(5), 2022, 899–916. <https://doi.org/10.1177/09622802211072512>.
- [23] Stilgoe, Jack, Richard Owen, and Phil Macnaghten. “Developing a Framework for Responsible Innovation.” *Research Policy* 42(9), 2013, 1568–1580. <https://doi.org/10.1016/j.respol.2013.05.008>.
- [24] Jafari, Akbar Anbar, Cagri Ozcinar, and Gholamreza Anbarjafari. “Responsible AI: The Good, The Bad, The AI.” arXiv preprint arXiv:2601.21095, 2026.